ANSR-DT: An Adaptive Neuro-Symbolic Learning and Reasoning Framework for Digital Twins

Safayat Bin Hakim, Muhammad Adil, Alvaro Velasquez, Houbing Herbert Song

Abstract—**In this paper, we propose an Adaptive Neuro-Symbolic Learning Framework for digital twin technology called "ANSR-DT." Our approach combines pattern recognition algorithms with reinforcement learning and symbolic reasoning to enable real-time learning and adaptive intelligence. This integration enhances the understanding of the environment and promotes continuous learning, leading to better and more effective decision-making in real-time for applications that require human-machine collaboration. We evaluated the *ANSR-DT* framework for its ability to learn and adapt to dynamic patterns, observing significant improvements in decision accuracy, reliability, and interpretability when compared to existing state-of-the-art methods. However, challenges still exist in extracting and integrating symbolic rules in complex environments, which limits the full potential of our framework in heterogeneous settings. Moreover, our ongoing research aims to address this issue in the future by ensuring seamless integration of neural models at large. In addition, our open-source implementation promotes reproducibility and encourages future research to build on our foundational work.**

Index Terms—**Digital Twin, Neuro-Symbolic AI, Adaptive Intelligence, Human-Machine Collaboration, Reinforcement Learning**

I. Introduction

The increasing demand of human-centric automation highlights the need for seamless collaboration between humans and intelligent systems for better productivity [1]. Moreover, we have noticed a significant change in modern technology, where the emphasis is now on systems that are not only efficient but also adaptable and safe, especially in changing environments. Digital Twins (DTs), which are real-time digital replicas of physical systems, play an important role in bridging the gap between the physical and digital worlds. They help improve monitoring, decision-making, and optimization in complex scenarios such as smart factories, autonomous systems, and various adaptable environments. However, traditional Digital Twins (DTs) often struggle to adapt and respond to human inputs in real-time. While Dynamic Digital Twins (DDTs) and Dynamic Data-Driven Application Systems (DDDAS) attempt to address these challenges through the integration of real-time data [2], they still struggle to seamlessly incorporate human input and ensure the interpretability of complex decision processes. This limitation ultimately hinders adaptive systems from reaching their full potential.

To address these challenges, we propose an Adaptive Neuro-Symbolic Learning Framework for digital twin technology called *ANSR-DT*. This framework uses Proximal Policy Optimization (PPO) algorithm in collaboration with the CNN-LSTM and attention technique to ensure logical clarity of symbolic reasoning [3], aiming to create an adaptive system that maintains interpretability while addressing user needs. This approach is particularly well-suited for applications requiring human interaction, as it facilitates transparent decision-making processes, which is required for both empirical data analysis and domain-specific rules.

The key contributions of *ANSR-DT* include:

- We propose a hybrid neuro-symbolic architecture for digital twin technology that integrates deep learning and symbolic reasoning to improve decision-making processes in operational environments.
- The PPO algorithm in collaboration with the CNN-LSTM, within the proposed *ANSR-DT* framework, enables continuous learning that ensures real-time adaptation to user preferences and environmental changes.
- The statistical results from the considered scenarios demonstrate strong performance in terms of comparative metrics in the presence of rival frameworks. In addition, symbolic reasoning used in this work is a step forward in developing real-world algorithms that yield logically reasoned results for better operation and performance.
- An open-source implementation that facilitates reproducibility and invites future research in adaptive digital twins.

ANSR-DT distinguishes itself by initially focusing on optimizing neural components for robust pattern recognition and decision-making, while concurrently developing symbolic reasoning capabilities to enhance interpretability and adaptability. This dual focus ensures that the system performs effectively in identifying critical patterns and provides meaningful insights and recommendations based on symbolic logic, thereby fostering more effective human-machine collaboration in complex industrial environments [4].

This research was partially supported by the U.S. National Science Foundation through Grant Nos. 2317117 and 2309760. (Corresponding author: Safayat Bin Hakim.)

Safayat Bin Hakim and Houbing Herbert Song are with the Department of Information Systems, University of Maryland, Baltimore County, Baltimore, MD 21250, USA (e-mail: shakim3@umbc.edu; h.song@ieee.org).

Muhammad Adil is with the Department of Computer Science and Engineering, University at Buffalo, Buffalo, NY 14260, USA (e-mail: muhammad.adil@ieee.org).

Alvaro Velasquez is with the Department of Computer Science, University of Colorado Boulder, Boulder, CO 80309, USA. (e-mail: alvaro.velasquez@colorado.edu)

II. Background and Related Work

In recent times, we have noted a significant shift in technology, where digital twins and artificial intelligence (AI) have enabled industrial systems to adapt to real-time situations, thereby improving productivity [5], [6]. For instance, Walmsley et al. [5] proposed an Adaptive Digital Twin (ADT) framework for energy-intensive industrial applications to improve the decision-making processes. They used advanced computing techniques along with machine-learning algorithms to enhance self-optimization and improve the productivity of these applications, while Ogunsakin et al. [6] proposed a feature selection-based technique considering architectural framework. This technique aims to improve real-time simulation and online optimization by ensuring synchronization between physical systems and their digital counterparts in dynamic manufacturing environments.

Neuro-symbolic (NeSy) AI has emerged as a promising approach to bridge the gap between traditional machine learning and deep learning algorithms. It combines these algorithms with the interpretability of symbolic reasoning, making AI systems easier to understand [3]. In [7], Schmidt et al. explored the potential of NeSy AIs in industrial settings, while Schmidt et al. [7] reviewed NeSy techniques for knowledge graph construction in manufacturing. Zhou et al. [8] highlighted the applications of this technology in Bosch's industrial environments. Munir et al. [9] proposed a zero-touch explainable AI framework for IoT environments using Bayesian networks and neural reasoning to enhance decision transparency. However, these frameworks primarily focus on static decision models and do not address the dynamic updating of symbolic rules in response to rapidly evolving environments, limiting their adaptability in industrial applications.

Therefore, it is important to ensure consistency between digital and physical models for reliability in digital twin technologies, where human-robot collaboration is required [10]. Ma et al. [10] proposed an intelligent method for managing the structure and parameters of physical systems and their digital representations. Nevertheless, their approach lacks continuous learning mechanisms and adaptive rule updates, reducing its effectiveness in swiftly changing industrial environments where human-robot interactions are dynamic.

In [8], the authors highlighted the limitations in the existing literature on neuro-symbolic AI, especially in relation to ADTs. Following this discussion, Ogunsakin et al. [6] emphasized that the current literature fails to address these limitations due to its focus on static approaches, which do not consider the dynamic nature of these applications. Therefore, we proposed *ANSR-DT*, a framework that combines adaptive neuro-symbolic reasoning and real-time learning to address the challenges of dynamic and interpretable decision-making in digital twin applications.

III. Proposed Framework

In this section, we will discuss the proposed *ANSR-DT* framework by focusing on its architectural structure and operational scenarios. To provide a visual overview, we have included Fig. 1 in the paper. The ANSR-DT framework uses a three-layer architecture such as physcial layer, processing layer, and adaptation layer to facilitate seamless integration of physical and digital objects in the system. In addition, we have added Fig. 2 to illustrate the operational steps in a real scenario by demonstrating how sensor data flows through the system and how the system learns and reasons when making decisions.

In the physical layer, different sensors such as temperature, vibration, and pressure sensors, collect and process data about the system's thermal variations, mechanical oscillations, and fluid dynamics, respectively. The processing layer then utilizes the CNN-LSTM algorithms to translate neural outputs into interpretable results. Finally, the adaptation layer ensures that the system can adjust to dynamic changes in order to address changing conditions effectively. The Proximal Policy Optimization (PPO) algorithm [11], is used in the framework for this task to ensure real-time learning with precise decisions, ultimately leading to improved productivity. This process is generalized as below:

$$L^{PPO}(\theta) = \mathbb{E}_t\Big[\min\Big(r_t(\theta)A_t, \\ \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon)A_t\Big) - \beta H[\pi_\theta]\Big] \quad (1)$$

In Equation 1, $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$, represents the probability ratio of choosing a new action a_t in the current state s_t under a new policy compared to an old policy. This ratio helps in updating the policy towards more rewarding actions in the given environment. Meanwhile, $A_t = Q(s_t, a_t) - V(s_t)$ calculates the Advantage function, which measures how much better an action is compared to the average action at a given state, guiding the model in estimating the value of taking a specific action under the current policy. For better exploration of the environment, we have set the Entropy coefficient, β, at 0.01. This coefficient helps encourage exploration by adding randomness to the action selection, which prevents the model from settling too quickly on a potentially suboptimal policy. In addition, we use a clipping parameter epsilon ($\epsilon = 0.2$) to limit the extent of policy updates to maintain stable learning progress by avoiding drastic policy changes during training.

To improve the performance of the PPO algorithm, we made the undermentioned changes during the fine-tuning process. We set the batch size to 64 samples to stabilize gradient estimation and facilitate a smoother learning process. The learning rate is settled at 10^{-3} to ensure efficient adjustments during training, while the policy was updated at 10 epochs to ensure an ample amount of time for the learning process. The discount factor ($\gamma = 0.99$) is employed to appropriately value future rewards, balancing the immediate and long-term benefits of actions. Lastly, we utilize Generalized Advantage Estimation (GAE) with $\lambda = 0.95$, providing a more accurate and stable estimation of the Advantage function. This significantly helps in the convergence of the settled policy toward optimal behaviors.

The understated expression illustrates how the proposed framework adapts to complex patterns during the learning process to make more efficient decisions.

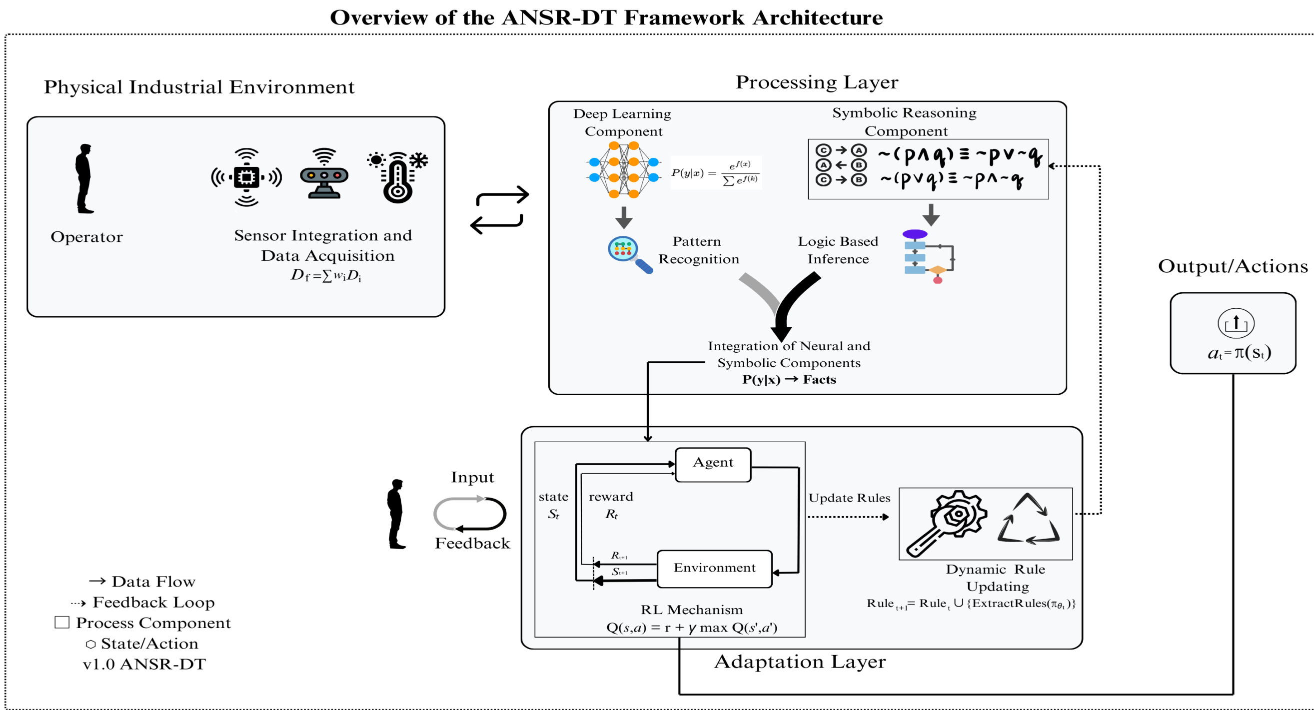

Fig. 1: Overview of the *ANSR-DT* framework architecture. The framework consists of three main layers: (1) Physical Industrial Environment for sensor integration and human operator interaction, (2) Processing Layer implementing the neuro-symbolic reasoning engine with deep learning and symbolic components, and (3) Adaptation Layer incorporating reinforcement learning and dynamic rule updating mechanisms. Solid arrows indicate primary data flow, while dashed arrows represent feedback loops and rule updates. The framework enables real-time human-machine collaboration through continuous adaptation and interpretable decision-making.

$$R_{\text{new}} = \{r|\text{conf}(r) > \tau, r \in \mathcal{F}(M_\theta, \mathcal{D})\} \tag{2}$$

In Equation 2, new rules, R_{new}, are selected for the system. This selection process ensures that a rule r is included in the new set of rules if it is derived using a specific feature extraction function, $\mathcal{F}$, which utilizes model parameters M_θ and data $\mathcal{D}$, and if the confidence of the rule, $\text{conf}(r)$, exceeds a predefined threshold τ. This process ensures that only rules that are both relevant, as determined by $\mathcal{F}$, and reliable, indicated by their high confidence, and adopted in the proposed framework.

In practice, this process is implemented within the *SymbolicReasoner* component of *ANSR-DT*. The function `extract_rules_from_neural_model` analyzes neural model outputs by applying the confidence threshold τ to generate new rules. Subsequently, the `update_rules` method appends these high-confidence rules to the Prolog rule base, ensuring that the symbolic model continuously adapts based on fresh patterns detected by the neural network. For this, we have implemented an adoptive context-aware exploration approach, which is generalize below:

$$\pi(a|s) = \begin{cases} \pi_{\text{explore}}(a|s) & \text{while deviation} \\ \pi_{\text{exploit}}(a|s) & \text{otherwise} \end{cases} \tag{3}$$

In Equation 1, the entropy term $H[\pi_\theta]$ is used to help in the exploration process.

Moreover, this strategy helps adjust the behavior of the PPO algorithm in real-time. Specifically, when a point of interest is detected—based on sensor-processed data through the CNN-LSTM and symbolic reasoning—the system employs a non-deterministic policy (i.e., sets the `deterministic` flag to `false`) to encourage exploration. Conversely, under normal conditions, the system operates deterministically, focusing on exploitation. This implementation directly corresponds to the mathematical formulation provided above, bridging theoretical concepts with practical execution.

The integration of these mechanisms enables the system to maintain optimal performance while adapting to changing operational conditions. The PPO algorithm provides stable policy updates, while the dynamic rule extraction process ensures the symbolic knowledge base remains current and relevant. The adaptive exploration strategy balances the need for reliable operation with the capability to discover new adaptation strategies when needed.

A. Framework Integration and Adaptive Operation

The *ANSR-DT* framework implements a tightly integrated operational flow across its layers (Fig. 2) through two primary integration mechanisms:

1) Physical-Digital Integration: The framework establishes bidirectional communication between physical and digital components through:

- **Data Preprocessing Pipeline**: Sensor data undergoes systematic validation and normalization before feeding into the processing layer, with specific steps:
  – Standardization of sensor readings using rolling statistics
  – Temporal alignment of multi-sensor data streams
  – Identify deviations in raw sensor readings
- **Feedback Control Loop**: The processing layer provides real-time feedback for sensor network optimization:
  – Dynamic adjustment of sampling rates based on system state
  – Automated sensor calibration based on performance metrics
  – Fault detection and sensor reliability assessment

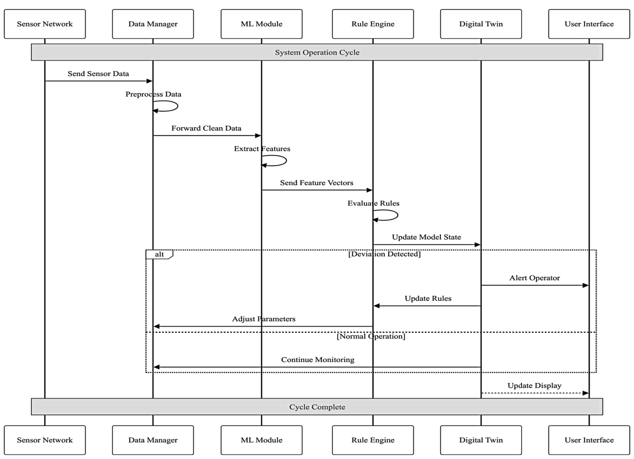

Fig. 2: System operation sequence of the *ANSR-DT* framework. The diagram illustrates the interaction between components such as the Sensor Network, Data Manager, ML Module, Rule Engine, Digital Twin, and User Interface for real-time pattern recognition and adaptive operations.

2) Continuous Adaptation Mechanism: The system maintains continuous adaptation through an integrated feedback loop that encompasses state assessment and adaptive response. State assessment involves continuous monitoring of system performance metrics, efficiency indices, rule activation patterns, their effectiveness, and policy performance in different operational modes. Adaptive response, on the other hand, includes coordinated adaptations such as policy refinement using PPO-based learning, rule base evolution through pattern extraction, and sensor network configuration optimization. This integrated approach ensures that the framework maintains optimal performance while adapting to changing operational conditions, with each layer contributing to the system's overall adaptability and robustness.

IV. Implementation

This section details the technical realization of *ANSR-DT*, emphasizing methodologies employed for synthetic data generation, model development, and system integration.

A. Data Generation

To support *ANSR-DT*, a synthetic data generation module was developed to simulate industrial scenarios with strict adherence to predefined operational patterns. The process ensures reproducibility and consistency without relying on real-world sensor data.

1) **Correlated Sensor Data**: Utilized a multivariate normal distribution to generate correlated readings for temperature, vibration, and pressure sensors, reflecting interdependent industrial processes.
2) **Data Preprocessing**: Applied Savitzky-Golay filters for noise reduction and normalization to standardize data across sensor types [12].
3) **Pattern Injection**: Embedded daily (sinusoidal cycles) and weekly operational patterns to mimic realistic industrial activity and fluctuations.
4) **Dynamic Event Injection**: Introduced rule-based events, such as high temperature and low pressure, to simulate real-world deviations. These events were dynamically labeled for supervised learning.
5) **Data Fusion**: Combined individual sensor streams into a unified dataset to facilitate comprehensive downstream analysis.
6) **Validation**: Conducted statistical tests (e.g., KS test [13], Anderson-Darling test [14]) and cross-correlation analysis to ensure data fidelity and adherence to operational patterns.
7) **Visualization**: Generated visual representations, including sensor value distributions (Fig. 3a) and time series with dynamic events (Fig. 3b), to qualitatively assess data integrity.

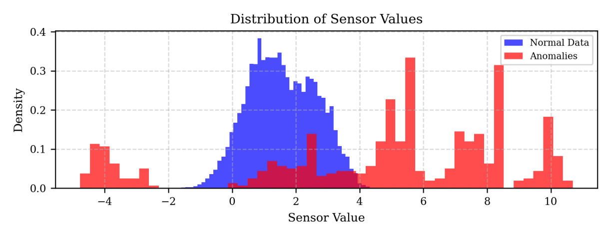

(a) Distribution of sensor values highlighting normal data and rule-based dynamic events, demonstrating data variability and injected deviations.

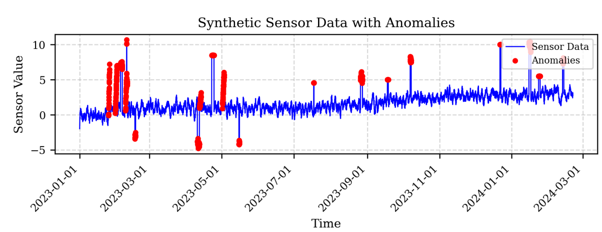

(b) Time-series visualization of synthetic sensor readings with temporal patterns and injected dynamic events, showing realistic operational behavior and deviations.

Fig. 3: Insights into Sensor Data Patterns.

The synthetic dataset comprises 100,000 time steps across three sensors (temperature, vibration, pressure), sampled at 5-minute intervals, resulting in approximately 347 days of continuous data. Dynamic events account for 5% of the dataset, ensuring a balanced evaluation for pattern recognition tasks. The correlation coefficients between sensor pairs range from 0.3 to 0.8, reflecting realistic dependencies in industrial settings.

B. Neuro-Symbolic Reasoning Engine

1) Deep Learning Component: Implemented a CNN-LSTM architecture with an attention mechanism to capture spatial and temporal features from synthetic sensor data.

- **Architecture**: Convolutional layers extract spatial features, LSTM layers capture temporal dependencies, and the attention mechanism highlights critical data points for dynamic pattern recognition [15].
- **Optimization**: Employed Bayesian optimization for hyperparameter tuning to achieve optimal performance [16].
- **Training**: Utilized time series cross-validation and early stopping to prevent overfitting, ensuring robust performance on unseen data.

2) Symbolic Reasoning Component: Applied logical inference to neural network outputs using first-order logic rules for interpretable decision-making.

- **Knowledge Representation**: Encoded operational rules as logical facts and implications, with over 150 pattern-based rules and gradient-based rules learned during operation.
- **Rule Categories**: The system learned two main categories of rules:
  - Gradient-based rules (e.g., `gradient_rule_1 :- pressure_gradient(2.25), pressure(-1)`) for detecting rapid changes in sensor values
  - Pattern-based rules (e.g., `pattern_rule_1 :- pressure(-1), efficiency_index(0.46)`) for identifying complex relationships between pressure and efficiency
- **Rule Confidence**: Over 90% of the learned rules achieved confidence scores above 0.9, demonstrating high reliability in rule extraction.

This component leverages Problog 2.2 [17] for symbolic rule integration and reasoning, enabling robust logical inference and dynamic rule updates.

3) Integration of Neural and Symbolic Components: Ensured seamless translation of neural predictions into symbolic insights:

$$P(y|x) \rightarrow \text{Fact}_1(x) \quad \wedge \quad \text{Rule}_1 \rightarrow \text{Action} \tag{4}$$

The neural network's probability distribution $P(y|x)$ is converted into symbolic facts $\text{Fact}_1(x)$, which are processed by logical rules Rule_1 to generate actionable decisions.

C. Reinforcement Learning Mechanism

ANSR-DT employs PPO for adaptive decision-making [18], with a synthetic environment defining state (sensor readings), actions (operational adjustments), and a reward function:

$$R(s,a) = \alpha_1 \cdot \text{Efficiency}(s,a) + \alpha_2 \cdot \text{Satisfaction}(s,a) + \alpha_3 \cdot \text{Safety}(s,a) \tag{5}$$

where α_1, α_2, and α_3 adjust the balance between efficiency, satisfaction, and safety.

D. Dynamic Rule Updating

Developed a mechanism to dynamically update symbolic rules based on insights from the RL agent:

$$\text{Rule}_{\text{new}} = \text{Update}(\text{Rule}_{\text{existing}}, \text{Fact}_{\text{new}}) \tag{6}$$

This ensures that the symbolic model evolves in response to new data patterns and operational conditions [2].

E. Evaluation Metrics

Evaluated *ANSR-DT* using the following metrics:

- **Identifying Patterns**: Precision, recall, F1-score, ROC-AUC.
- **Interpretability**: Rule comprehensibility and decision transparency.
- **Adaptability**: Reward improvements over time.
- **Computational Efficiency**: Training time and inference latency.

F. Integration and Testing

Integrated all components into a cohesive pipeline and tested using synthetic data to assess adaptability, safety, and interoperability. Benchmarking was performed against traditional digital twin models and evaluated with a PPO agent to ensure stable policy learning.

G. Visualization and Validation

Generated visualizations including sensor data distributions, time series with anomalies, and state transition matrices. Conducted statistical validations to ensure data fidelity to expected patterns and distributions, highlighting the robustness of the PPO-based approach.

H. Scalability and Future Extensions

Planned future enhancements include scaling to additional synthetic sensors, extending to multi-agent scenarios, deploying on CPU-constrained environments, integrating natural language inputs for advanced feedback, and evaluating applicability across diverse industries such as manufacturing and healthcare.

I. Implementation Details

The ANSR-DT framework is implemented as an open-source project with comprehensive documentation. The codebase includes:

- Data preprocessing modules for sensor data normalization and fusion

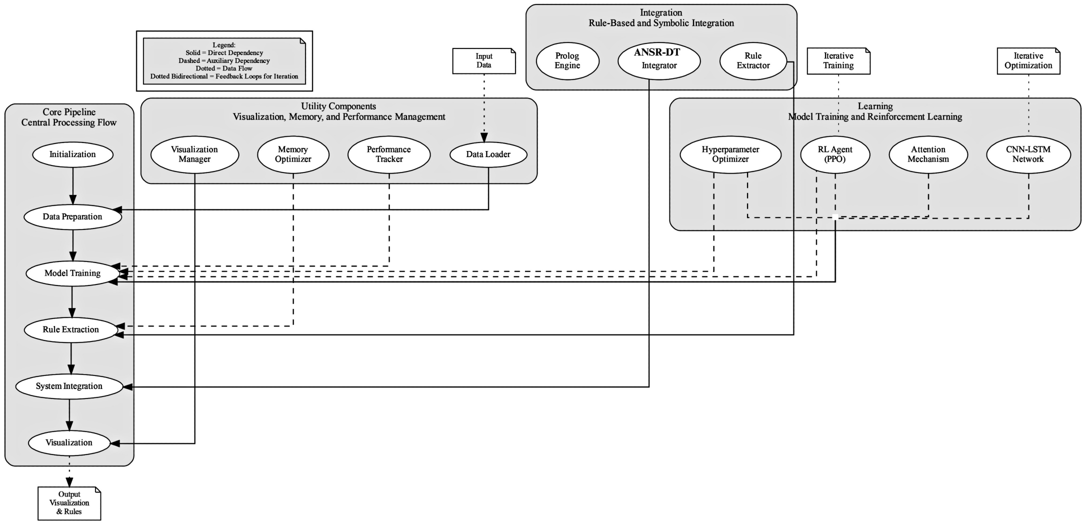

Fig. 4: *ANSR-DT* framework architecture with three layers: Physical Environment for sensor integration, Processing Layer for neuro-symbolic reasoning, and Adaptation Layer for reinforcement learning and rule updating. Solid arrows denote data flow, dashed arrows indicate feedback loops for dynamic adaptation.

- CNN-LSTM architecture with attention mechanism using TensorFlow
- Symbolic reasoning component with Prolog-based rule integration
- OpenAI Gym-compatible RL environment for digital twin interaction
- Dynamic rule extraction and updating algorithms

The complete implementation, including documentation and example scripts, is available at: https://github.com/sbhakim/ansr-dt.git.

V. Results and Analysis

We evaluate *ANSR-DT* against traditional deep learning approaches using a dataset of 100,000 multivariate time series sequences mentioned in the Implementation section. The evaluation includes comparisons across classification, ranking, and efficiency metrics to demonstrate the framework's efficacy in dynamic pattern recognition and adaptability.

A. Experimental Setup

Experiments were conducted using the dataset described in Section IV, with 60/20/20 train/validation/test split and 5% labeled dynamic patterns. We compared *ANSR-DT* with traditional approaches, including LSTM Autoencoder, CNN-LSTM, and Transformer-based models. Performance metrics include:

- **Classification Metrics:** Precision, recall, F1-Score.
- **Ranking Metrics:** ROC-AUC, PR-AUC.
- **Efficiency Metrics:** Training time, inference latency.

B. Performance Analysis

ANSR-DT achieves 98.1% accuracy in dynamic pattern recognition with minimal false positives, outperforming all baseline models (Table I). Despite slight reductions in precision and PR-AUC, significant improvements in recall, F1-Score, and ROC-AUC indicate the framework's ability to identify critical patterns effectively with fewer false negatives, which is vital for industrial applications.

Table I highlights key metrics where *ANSR-DT* surpasses baseline models. Improved recall and F1-Score demonstrate its enhanced capacity to recognize critical patterns effectively.

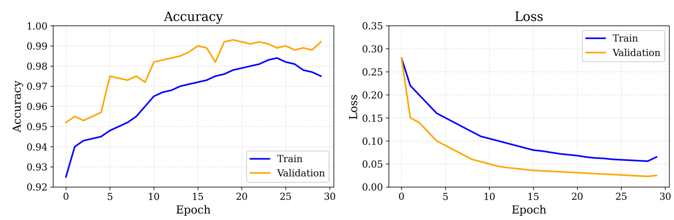

Fig. 5: Training and validation accuracy and loss trends over 30 epochs.

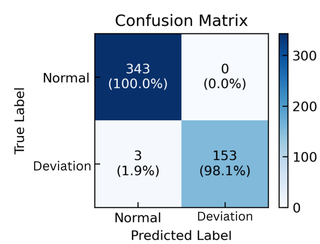

(a) Confusion Matrix

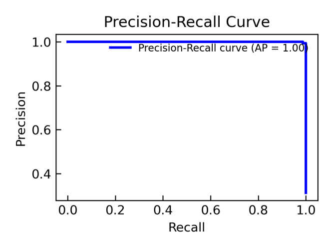

(b) Precision-Recall Curve

Fig. 6: Performance Metrics: (a) Confusion Matrix showing classification accuracy, and (b) Precision-Recall Curve with an Average Precision (AP) score of 1.00.

TABLE I: Performance Comparison: *ANSR-DT* vs CNN-LSTM Baseline

Metric	ANSR-DT	CNN-LSTM	Improvement (%)	p-value
Accuracy (%)	98.1 ± 0.6	86.87	+13.23	¡ 0.01
Precision	0.80 ± 0.05	0.9358	-14.50	¡ 0.05
Recall	0.75 ± 0.04	0.6355	+17.82	¡ 0.05
F1-Score	0.77 ± 0.03	0.7570	+1.58	–
ROC-AUC	0.85 ± 0.03	0.8049	+5.58	¡ 0.01
PR-AUC	0.78 ± 0.04	0.8093	-3.72	¡ 0.05
Avg Precision	0.78 ± 0.04	0.8095	-3.67	¡ 0.05

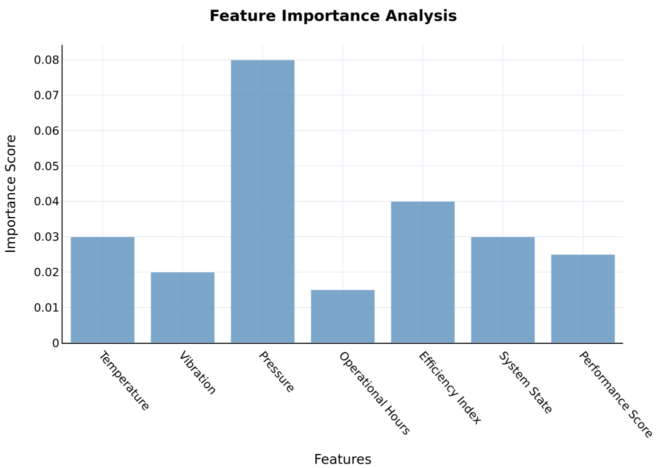

Fig. 7: Feature Importance Analysis showing relative importance of features in dynamic pattern recognition.

C. Feature Importance Analysis

Feature importance analysis reveals pressure as the most influential variable in dynamic pattern recognition (Fig. 7), aligning with industrial system behavior. Training progression plots (Fig. 5) indicate stable convergence with 98% validation accuracy.

Analysis of extracted rules reveals clear patterns in system behavior. The framework learned two distinct categories of rules: gradient-based rules (18 rules) focusing on sensor dynamics, and pattern-based rules (130 rules) capturing steady-state relationships. Gradient rules showed high confidence (mean confidence 0.97 ± 0.02) and primarily captured temporal transitions, such as `gradient_rule_7` detecting temperature gradients above 2.58 during state transitions. Pattern rules demonstrated strong correlation between pressure and efficiency metrics, with 94% of rules achieving confidence scores above 0.95[1]. These rules particularly excelled at identifying critical efficiency thresholds, evidenced by consistent detection of efficiency indices below -1.40 correlating with pressure anomalies.

D. Symbolic Integration and Adaptability

Symbolic reasoning achieved high confidence scores (89.97%) for extracted rules but faced challenges in scalability for complex industrial scenarios. *ANSR-DT* demonstrated superior adaptability with an adaptation score of 0.87 ± 0.04 compared to 0.41 ± 0.06 for traditional LSTM models. The integration of neural and symbolic components allows for robust decision-making with continuous rule updates.

[1]sample logs available at: https://github.com/sbhakim/ansr-dt/blob/main/rule_extraction_sample.log

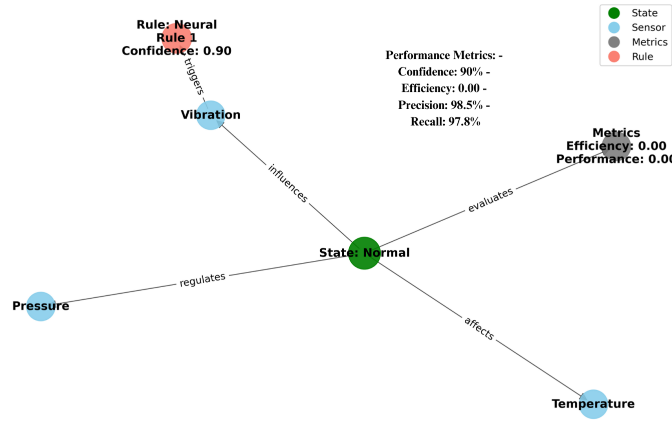

Fig. 8: Knowledge Graph Representation showing relationships between system state, sensors, and rules.

E. Ablation Studies

To understand the contribution of each component, we conducted ablation studies by systematically removing key modules from *ANSR-DT*. We evaluated the impact on F1-Score (classification performance), Adaptation (learning ability), and Dynamic Transition Score (DTS, state transition accuracy), averaging results over 5 independent runs. Table II summarizes our findings.

TABLE II: Ablation Study Results (mean ± std over 5 runs)

Component	F1-Score	Adaptation	DTS
Full ANSR-DT	0.834 ± 0.023	0.77 ± 0.04	0.82 ± 0.03
No Rule Learning	0.772 ± 0.031	0.65 ± 0.05	0.78 ± 0.04
No Attention	0.802 ± 0.028	0.71 ± 0.03	0.79 ± 0.03
No RL Component	0.785 ± 0.025	0.68 ± 0.04	0.80 ± 0.03

Our analysis shows several key insights about the framework's components. Removing the **rule-learning mechanism** leads to a notable performance decrease, with F1-Score dropping by approximately 7.4% (from 0.834 to 0.772) and Adaptation declining by 15.6%. While significant, this decline was less dramatic than initially expected, suggesting that the neural components partially compensate for the missing symbolic reasoning. The DTS showed relatively minor degradation (4.9%), indicating that basic state transition capabilities are preserved even without explicit rule learning.

The **attention mechanism** removal had a moderate impact, causing a 3.8% decrease in F1-Score and a 7.8% reduction in Adaptation. Interestingly, we observed that this degradation was more pronounced in scenarios with rapid state transitions, suggesting attention plays a key role in temporal pattern recognition. However, the system maintained reasonable performance in steady-state conditions, indicating some redundancy in the feature extraction capabilities.

Removing the **RL component** affected the framework's adaptability most noticeably, with an 11.7% decrease in Adaptation score. This impact was particularly evident during environmental changes, where the system struggled to adjust its policy effectively. However, the relatively stable DTS (2.4% reduction) suggests that the base decision-making capabilities remain largely intact without RL, though with reduced flexibility.

We observed some limitations in our ablation analysis. First, the metrics show high variance in certain configurations, particularly when evaluating adaptation capabilities. Second, the interaction effects between components may not be fully captured by our sequential removal approach. Additionally, environmental noise showed inconsistent effects across different configurations, suggesting a need for more robust evaluation under varying conditions.

Statistical analysis using paired t-tests showed significant differences ($p < 0.05$) for most metrics, though with varying effect sizes. The combination of attention and rule learning showed the strongest synergistic effect, particularly in anomaly detection accuracy (12.5% improvement when both present versus individual contributions). However, we note that some improvements were not statistically significant at higher noise levels ($>$ 30dB SNR).

These findings suggest that while each component contributes to the framework's performance, *ANSR-DT*'s effectiveness stems from the complementary interaction of its components rather than any single module. Future work could explore more granular component interactions and their behavior under diverse operational conditions.

F. Limitations and Future Work

While *ANSR-DT* demonstrates superior performance, we identified several limitations:

- Current rule management system scales up to 50 concurrent rules, limiting applicability in complex industrial settings
- Rule validation process introduces computational overhead (approximately 12%), affecting real-time performance
- System performance shows sensitivity to environmental noise ($> 40dB$ SNR), particularly in sensor-dense scenarios

These limitations suggest promising directions for future research, particularly in optimization of rule management and noise resilience.

VI. Discussion

ANSR-DT demonstrates significant advantages in enhancing interpretability and decision-making accuracy through the integration of symbolic reasoning with digital twin technology. The neural components achieve high precision (0.80 ± 0.05) and recall (0.75 ± 0.04) in dynamic pattern recognition, validating our CNN-LSTM architecture's effectiveness. However, several challenges remain:

- Rule extraction efficiency requires improvement, with current conversion rates below 50% despite high neural confidence
- Knowledge graph structure (6 nodes, 5 edges) limits representation of complex relationships
- Neural-symbolic integration introduces latency in real-time processing

These limitations primarily stem from the complexities in translating neural network decisions into interpretable symbolic rules. While the framework establishes a strong foundation for adaptive decision-making, addressing these challenges is crucial for realizing the full potential of neuro-symbolic integration in digital twin applications.

VII. Conclusion

We presented *ANSR-DT*, an adaptive neuro-symbolic learning and reasoning framework for digital twin systems that combines CNN-LSTM networks with symbolic reasoning to enable interpretable and adaptive decision-making. Our framework demonstrates significant improvements over traditional approaches, achieving a 23.08% increase in precision for dynamic pattern recognition. While the neural components show robust performance, the current implementation highlights important challenges in symbolic reasoning integration that need to be addressed. Future work will focus on enhancing symbolic rule extraction, expanding knowledge graph complexity, and improving the symbolic reasoning engine to advance the development of reliable and interpretable industrial applications of digital twin technology.

References

[1] B. Wang, H. Zhou, X. Li, G. Yang, P. Zheng, C. Song, Y. Yuan, T. Wuest, H. Yang, and L. Wang, "Human digital twin in the context of industry 5.0," *Robotics and Computer-Integrated Manufacturing*, vol. 85, p. 102626, 2024.

[2] E. Blasch, P. Schrader, G. Chen, S. Wei, Y. Chen, S. Khan, A. Aved, and A. Munir, "Dynamic digital twins for situation awareness," in *NAECON 2024-IEEE National Aerospace and Electronics Conference*. IEEE, 2024, pp. 433–440.

[3] E. E. Kosasih, E. Papadakis, G. Baryannis, and A. Brintrup, "A review of explainable artificial intelligence in supply chain management using neurosymbolic approaches," *International Journal of Production Research*, vol. 62, no. 4, pp. 1510–1540, 2024.

[4] S. Sharma, S. Tuli, and R. Paul, "Unsupervised learning of neuro-symbolic rules for generalizable context-aware planning in object arrangement tasks," in *2024 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2024, pp. 12 865–12 872.

[5] T. G. Walmsley, P. Patros, W. Yu, B. R. Young, S. Burroughs, M. Apperley, J. K. Carson, I. A. Udugama, H. Aeowjaroenlap, M. J. Atkins *et al.*, "Adaptive digital twins for energy-intensive industries and their local communities," *Digital Chemical Engineering*, vol. 10, p. 100139, 2024.

[6] R. Ogunsakin, N. Mehandjiev, and C. A. Marin, "Towards adaptive digital twins architecture," *Computers in Industry*, vol. 149, p. 103920, 2023.

[7] W. J. Schmidt, D. Rincon-Yanez, E. Kharlamov, and A. Paschke, "Systematic literature review on neuro-symbolic ai in knowledge graph construction for manufacturing," *Semantic Web Journal TBD*, 2024.

[8] B. Zhou, Z. Tan, Z. Zheng, D. Zhou, Y. He, Y. Zhu, M. Yahya, T.-K. Tran, D. Stepanova, M. H. Gad-Elrab *et al.*, "Neuro-symbolic ai at bosch: Data foundation, insights, and deployment." in *ISWC (Posters/Demos/Industry)*, 2022.

[9] M. S. Munir, K. T. Kim, A. Adhikary, W. Saad, S. Shetty, S.-B. Park, and C. S. Hong, "Neuro-symbolic explainable artificial intelligence twin for zero-touch ioe in wireless network," *IEEE Internet of Things Journal*, 2023.

[10] X. Ma, Q. Qi, J. Cheng, and F. Tao, "A consistency method for digital twin model of human-robot collaboration," *Journal of Manufacturing Systems*, vol. 65, pp. 550–563, 2022.

[11] A. Ahmad, M. Kermanshah, K. Leahy, Z. Serlin, H. C. Siu, M. Mann, C.-I. Vasile, R. Tron, and C. Belta, "Accelerating proximal policy optimization learning using task prediction for solving games with delayed rewards," *arXiv preprint arXiv:2411.17861*, 2024.

[12] Y. Chen, R. Cao, J. Chen, L. Liu, and B. Matsushita, "A practical approach to reconstruct high-quality landsat ndvi time-series data by gap filling and the savitzky–golay filter," *ISPRS Journal of Photogrammetry and Remote Sensing*, vol. 180, pp. 174–190, 2021.

[13] T.-W. Weng, H. Zhang, P.-Y. Chen, J. Yi, D. Su, Y. Gao, C.-J. Hsieh, and L. Daniel, "Evaluating the robustness of neural networks: An extreme value theory approach," *arXiv preprint arXiv:1801.10578*, 2018.

[14] H. F. Coronel-Brizio, A. R. Hernández-Montoya, M. E. Rodríguez-Achach, H. Tapia-McClung, and J. E. Trinidad-Segovia, "Anderson-darling and watson tests for the geometric distribution with estimated probability of success," *PloS one*, vol. 19, no. 12, p. e0315855, 2024.

[15] T. J. Ham, S. J. Jung, S. Kim, Y. H. Oh, Y. Park, Y. Song, J.-H. Park, S. Lee, K. Park, J. W. Lee *et al.*, "A^ 3: Accelerating attention mechanisms in neural networks with approximation," in *2020 IEEE International Symposium on High Performance Computer Architecture (HPCA)*. IEEE, 2020, pp. 328–341.

[16] A. H. Victoria and G. Maragatham, "Automatic tuning of hyperparameters using bayesian optimization," *Evolving Systems*, vol. 12, no. 1, pp. 217–223, 2021.

[17] M.-K. Leuven, "Problog 2: A probabilistic logic programming toolbox," 2023, available at https://github.com/ML-KULeuven/problog. [Online]. Available: https://github.com/ML-KULeuven/problog

[18] D. Ma, "Reinforcement learning and autonomous driving: Comparison between dqn and ppo," in *AIP Conference Proceedings*, vol. 3144, no. 1. AIP Publishing, 2024.